

Молодежная школа ННГУ

OpenVINO stories

Dmitry Kurtaev



intel®

Domains



Retail



Industrial



Medical



Xiaogu Technology

- Facial recognition payment solution
- Android x86 OS
- Intel Apollo Lake N3350
 - CPU: Intel Atom® x7-E3950
 - GPU: Intel® HD Graphics 505



Algorithm	Execution time after acceleration with OpenVINO	
	CPU	GPU
Face detection	97 ms	48 ms
Face quality detection	40 ms	23 ms
Face liveness recognition	3 ms	2.5 ms
Face recognition	112 ms	57 ms

White paper: <https://www.intel.cn/content/www/cn/zh/internet-of-things/xiaogu-3d-face-frictionless-payment-solution-en.html>

Xiaogu Technologies experience


- Do not implement classes in headers!
 - [RTTI/Exceptions Not Working Across Library Boundaries](#)
- cl_cache is powerful
- OpenCL for Android: Use Clang to build Clang to build Clang


Chinese License plate recognition

- 33 province letters
- regular, transit, public, military



 [license-plate-recognition-barrier-0001](https://github.com/license-plate-recognition-barrier-0001)





Weld porosity detection

- Robotic arc weld
- Human vision inspection leads to many false positives



 [/openvinotoolkit/open_model_zoo/blob/master/models/intel/weld-porosity-detection-0001](https://github.com/openvinotoolkit/open_model_zoo/blob/master/models/intel/weld-porosity-detection-0001)





Oil & Gas: fault segmentation





Low-Precision Inference for
High-Performance,
Fault Detection using
3D Seismic Data

- Accelerating fault detection in 3D Seismic data using OpenVINO - Reducing time to the first Oil
- Analyzing 3D Seismic Data using Intel Distribution of OpenVINO Toolkit


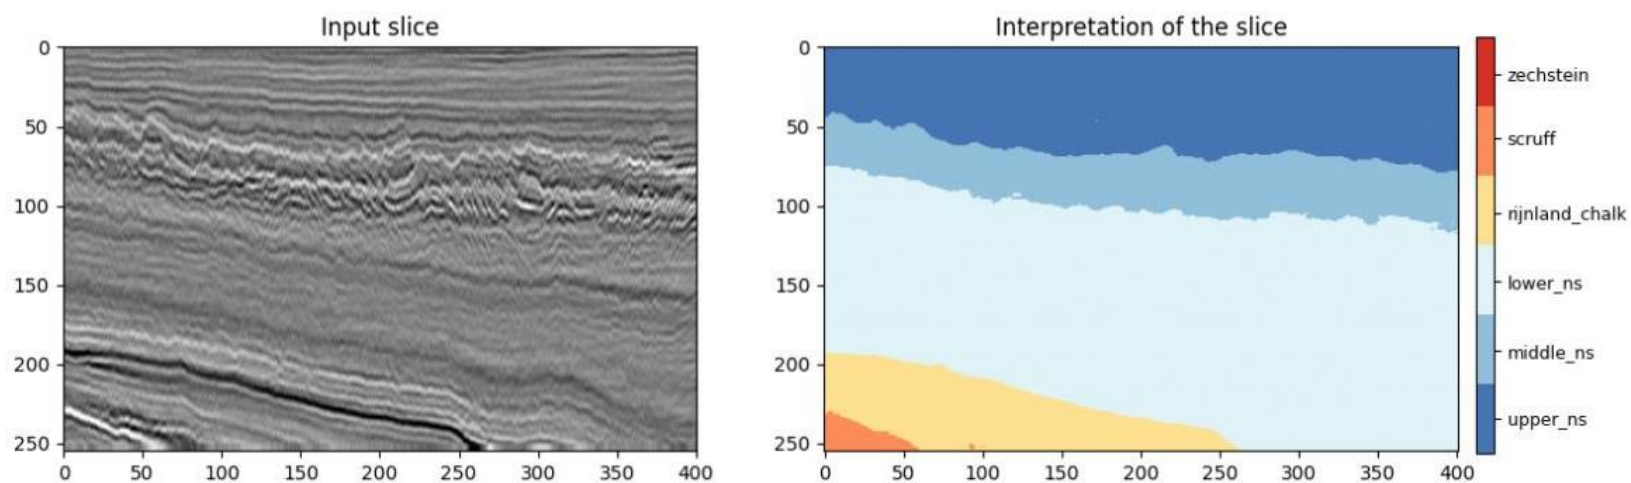
Oil & Gas: fault segmentation






Oil & Gas: salt detection

Facies classification results



Original 3D seismic section

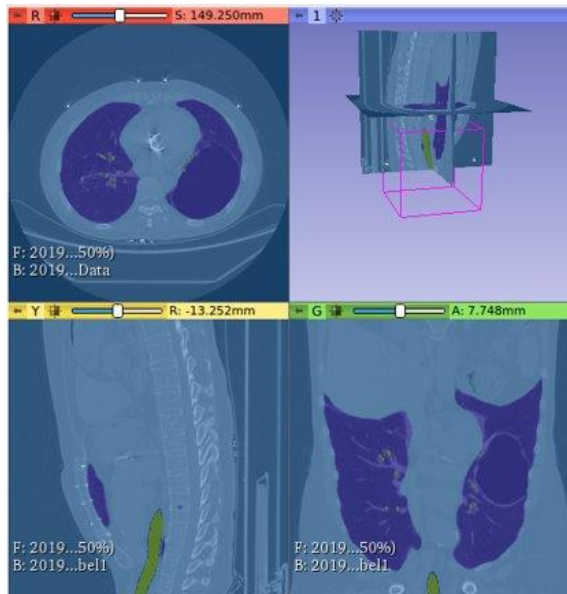


Salt prediction

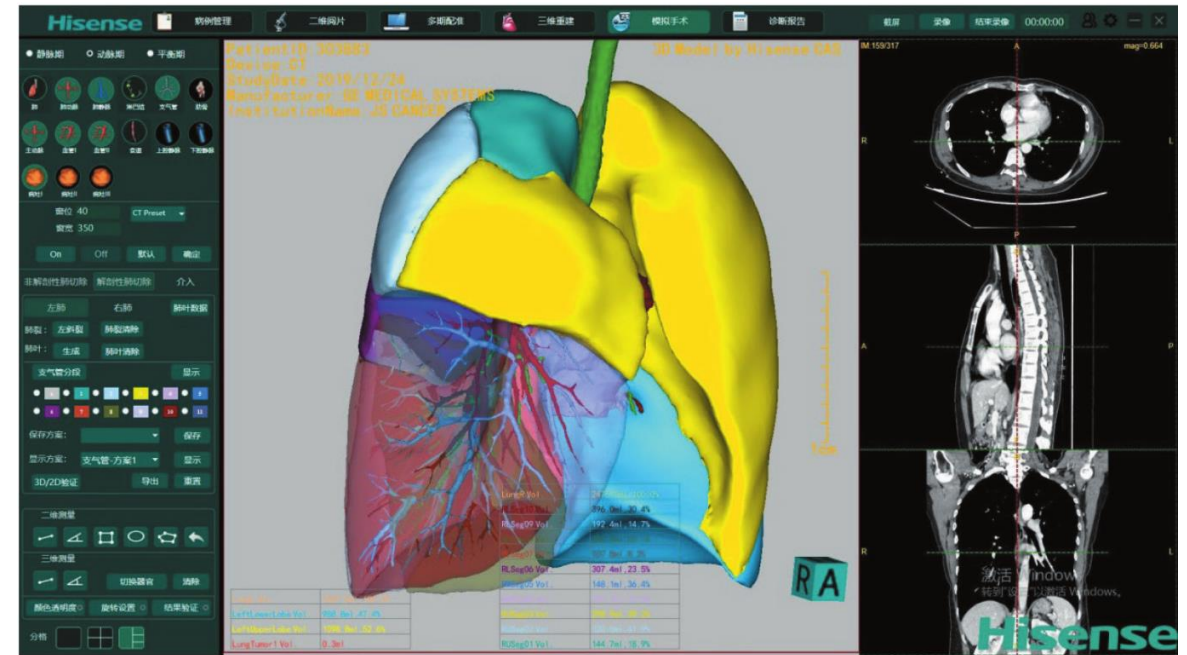



Surgery assistants

- Pulmonary surgery assistant
- NiftyNet framework based
- Dense VNet (3D UNet)



Hisense Medical



White paper: [Hisense Medical Computer-assisted Surgery \(CAS\) System for Pulmonary Surgery based on Intel® Architecture](#)


Genomics

Nucleotides

- * Adenine
- * Cytosine
- * Guanine
- * Thymine


Base pairs

A - T
C - G




Sequencing

- * WGS = whole genome sequencing



Genomics



<https://nanoporetech.com/resource-centre/how-nanopore-sequencing-works-animation>

 [/nanoporetech/bonito](https://github.com/nanoporetech/bonito)

- Long reads sequencing based on DL network (PyTorch)
- 1D signal processing

Genomics



MinION




GridION



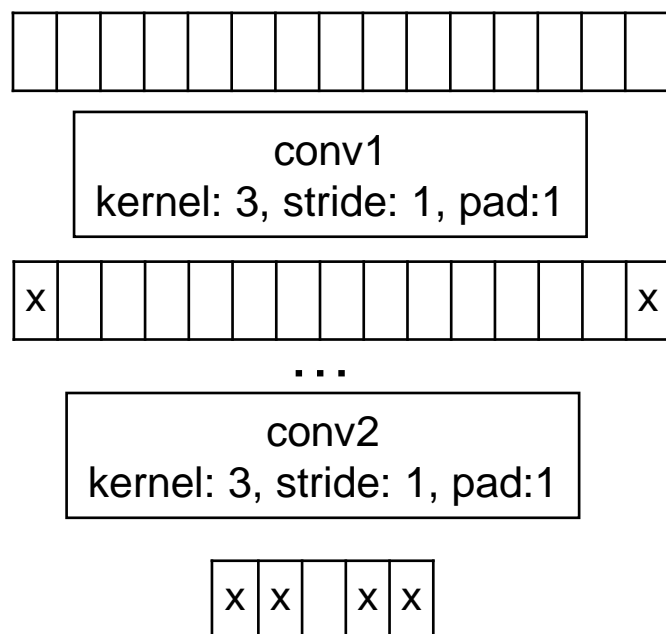
PromethION

Genomics

1. Full read processing – accurate but not scalable



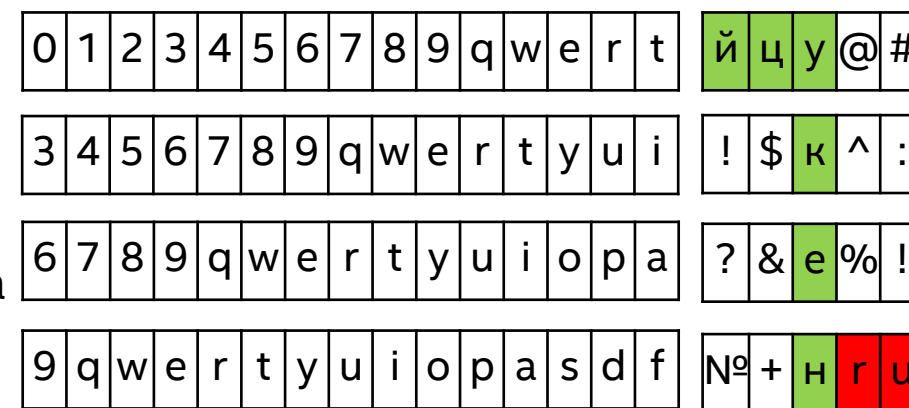
2. Chunks processing



Chunk size – any
overlap – estimate by “clean” area
(depends on the topology, not
chunk size)

Chunk size: 15

Overlap: 12



Bonito example: overlap: 2988

16154 input length is

- 22 chunks of 3600 numbers
- 5 chunks of 6000 numbers

Variant calling

SNPs (snips) - single nucleotide polymorphisms

INDEL – insertion or deletion of nucleotides

ACGTGAAGAATCATA
ACGTGATGAATCATA

- Genetic disorder
 - Monogenic (melanoma, phenylketonuria, mucopolysaccharidoses)
 - Polygenic (asthma, coronary heart disease, and diabetes)
- Only 1.5% of genes are “active” (exome)
- *de novo* sequencing: genome functioning rather than structure (i.e. viruses)


Google DeepVariant

- Inception v3 based model
- Docker as primary API

External Solutions

The following pipelines are not created or maintained by the [Genomics team in Google Health](#). Please contact the relevant teams if you have any questions or concerns.

Name	Description
Running DeepVariant on Google Cloud Platform	Docker-based pipelines optimized for cost and speed. Code can be found here .
DeepVariant-on-spark from ATGENOMIX	A germline short variant calling pipeline that runs DeepVariant on Apache Spark at scale with support for multi-GPU clusters (e.g. NVIDIA DGX-1).
NVIDIA Clara Parabricks	An accelerated DeepVariant pipeline with multi-GPU support that runs our WGS pipeline in just 40 minutes, at a cost of \$2-\$3 per sample. This provides a 7.5x speedup over a 64-core CPU-only machine at lower cost.
DNAnexus DeepVariant App	Offers parallelized execution with a GUI interface (requires platform account).
Nextflow Pipeline	Offers parallel processing of multiple BAMs and Docker support.
DNAstack Pipeline	Cost-optimized DeepVariant pipeline (requires platform account).



Input: 1x100x221x6

Output classification:

0 - homozygous reference allele (no change)


1 - heterozygous

2 - homozygous alternative allele

Compressed sense



raw signal (k-space)



visualization

Is that possible to sample less MRI / CT data?

- Faster to scan, healthier (physical/mental), lighter storage

Compressed sense







image of origin slice
(from kspace)



sampling mask
0.82 of zeros
(applied to kspace)



slice + mask as is
PSNR: 21.5




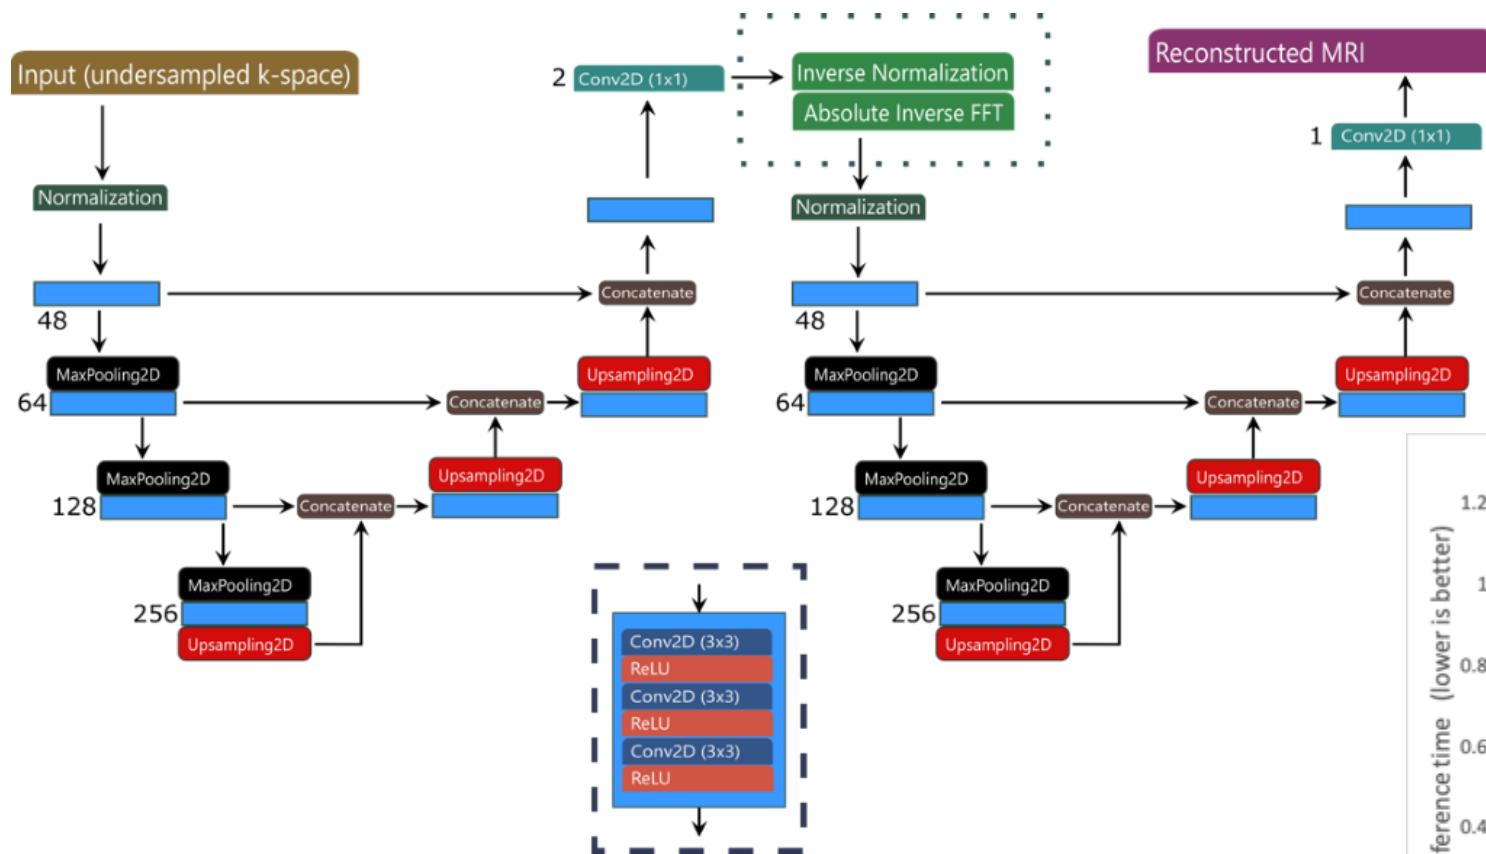
Reconstructed via network
PSNR: 34.8



project (MIT): <https://github.com/rmsouza01/Hybrid-CS-Model-MRI/>

dataset (Creative Commons Attribution-NoDerivatives 4.0 International Public License):
<https://sites.google.com/view/calgary-campinas-dataset/home>

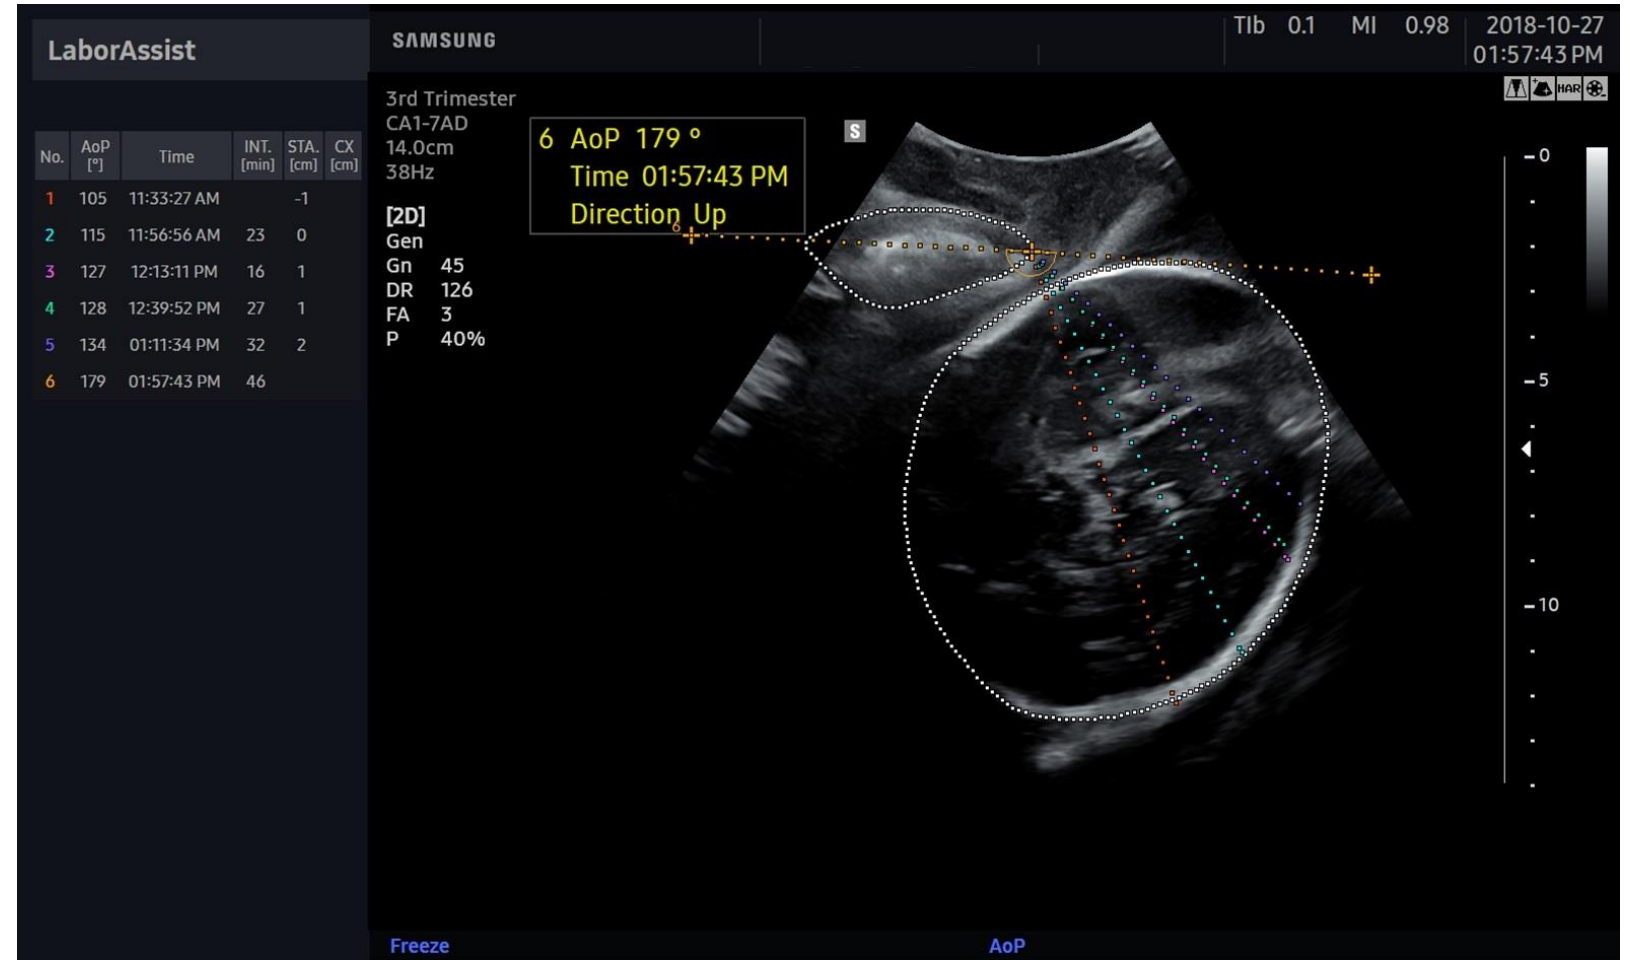
Compressed sense



White paper: [Philips Healthcare MRI Image Reconstruction](#)

Samsung: Fetal Ultrasound Workflow

- Reduce manual operations
- Powered by Intel Core i3



White paper: <https://www.intel.com/content/www/us/en/healthcare-it/samsung-medical-imaging-brief.html>

